

Fire up libraries

1 - Josemar Figueiredo Pereira

In [1]:

```
import matplotlib
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
%matplotlib inline
from sklearn import linear_model
```

Load some house sales data

Dataset is from house sales in King County, the region where the city of Seattle, WA is located.

In [2]:

```
sales = pd.read_csv('home_data.csv')
```


In [3]:

```
sales.head()
```

Out[3]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot
0	7129300520	20141013T000000	221900	3	1.00	1180	5650
1	6414100192	20141209T000000	538000	3	2.25	2570	7242
2	5631500400	20150225T000000	180000	2	1.00	770	10000
3	2487200875	20141209T000000	604000	4	3.00	1960	5000
4	1954400510	20150218T000000	510000	3	2.00	1680	8080

5 rows × 21 columns




In [4]:

```
sales[sales['id']==1839920160]
```

Out[4]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft
11860	1839920160	20140714T000000	432000	3	2.0	1870	708

1 rows × 21 columns



In [5]:

```
sales.keys()
```

Out[5]:

```
Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',  
      'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',  
      'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',  
      'lat', 'long', 'sqft_living15', 'sqft_lot15'],  
      dtype='object')
```

In [6]:

```
sales.shape
```

Out[6]:


```
(21613, 21)
```

Exploring the data for housing sales

The house price is correlated with the number of square feet of living space.

In [7]:

```
plt.grid('on')
plt.scatter(sales['sqft_living'], sales['price'])
plt.show()
```



Create a simple regression model of sqft_living to price

Split data into training and testing.

We use random_state=200 so that everyone running this notebook gets the same results. In practice, you may set a random seed.

In [8]:

```
train_data = sales.sample(frac=0.8, random_state=200)
test_data = sales.drop(train_data.index)
print(train_data.shape, test_data.shape)
```

```
(17290, 21) (4323, 21)
```


In [9]:

```
train_data.head()
```

Out[9]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft
11860	1839920160	20140714T000000	432000	3	2.00	1870	708
12446	6705850140	20141009T000000	750000	4	2.75	3170	763
10556	924069190	20140819T000000	440000	3	1.75	2000	118
4828	3211270170	20140523T000000	404000	4	3.00	4060	356
3502	9523103001	20141013T000000	389000	2	1.00	850	327

5 rows × 21 columns




In [10]:

```
test_data.head()
```

Out[10]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_living2
3	2487200875	20141209T000000	604000	4	3.0	1960	5000
4	1954400510	20150218T000000	510000	3	2.0	1680	8080
5	7237550310	20140512T000000	1225000	4	4.5	5420	10190
17	6865200140	20140529T000000	485000	4	1.0	1600	4300
18	16000397	20141205T000000	189000	2	1.0	1200	9850

5 rows × 21 columns



Build the regression model using only sqft_living as a feature

In [11]:

```
x_train = train_data['sqft_living'].values.reshape(-1,1)
y_train = train_data['price'].values.reshape(-1,1)
```

In [12]:


```
simple_model = linear_model.LinearRegression()
simple_model.fit(x_train, y_train)
```

Out[12]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

In [13]:

```
plt.grid('on')
plt.scatter(x_train, y_train)
plt.plot(x_train, simple_model.predict(x_train), color='red', linewidth=2)
plt.show()
```



Let's show what our predictions look like


In [14]:

```
x_test = test_data['sqft_living'].values.reshape(-1,1)
y_test = test_data['price'].values.reshape(-1,1)

y_pred = simple_model.predict(x_test)
```

In [15]:

```
plt.grid('on')
plt.scatter(x_test, y_test)
plt.plot(x_test, y_pred, color='red', linewidth=2)
plt.show()
```



Evaluate the simple model

In [16]:

```
def rmse(predictions, targets):
    return np.sqrt(((predictions - targets) ** 2).mean())
```

In [17]:

```
print('intercept:', simple_model.intercept_, 'coefficients:', simple_model.coef_)
# The mean squared error
print("RMSE: %.2f" % (rmse(y_pred, y_test)))
```

```
intercept: [-46493.04519733] coefficients: [[ 282.27187583]]
RMSE: 254323.39
```

RMSE of about \$254.323,39

Explore other features in the data

To build a more elaborate model, we will explore using more features.

In [18]:

```
my_features = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'zipcode']
```

In [19]:


```
sales[my_features].describe()
```

Out[19]:

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	
count	21613.000000	21613.000000	21613.000000	2.161300e+04	21613.000000	21613.000000
mean	3.370842	2.114757	2079.899736	1.510697e+04	1.494309	9807.000000
std	0.930062	0.770163	918.440897	4.142051e+04	0.539989	53.500000
min	0.000000	0.000000	290.000000	5.200000e+02	1.000000	9800.000000
25%	3.000000	1.750000	1427.000000	5.040000e+03	1.000000	9803.000000
50%	3.000000	2.250000	1910.000000	7.618000e+03	1.500000	9806.000000
75%	4.000000	2.500000	2550.000000	1.068800e+04	2.000000	9811.000000
max	33.000000	8.000000	13540.000000	1.651359e+06	3.500000	9819.000000

In [20]:

```
#sales.show(view='BoxWhisker Plot', x='zipcode', y='price')
plt.grid('on')
plt.scatter(sales['zipcode'], sales['price'])
plt.show()
```



98039 is the most expensive zip code.

Build a regression model with more features

In [21]:

```
#my_features_model = (train_data,target='price',features=my_features,validation_set=None)
x_train = train_data[my_features].values.reshape(-1,len(my_features))
y_train = train_data['price'].values.reshape(-1,1)
```

In [22]:

```
mult_model = linear_model.LinearRegression()
mult_model.fit(x_train, y_train)
```

Out[22]:

```
LinearRegression(copy_X=True, fit_intercept=True, n_jobs=1, normalize=False)
```

Comparing the results of the simple model with adding more features

In [23]:

```
x_test = test_data[my_features].values.reshape(-1,len(my_features))
y_test = test_data['price'].values.reshape(-1,1)

y_pred = mult_model.predict(x_test)
```

In [24]:

```
print('intercept:', mult_model.intercept_, 'coefficients:', mult_model.coef_)
# The mean squared error
print("RMSE: %.2f" % (rmse(y_pred, y_test)))
```

```
intercept: [-56348418.94956037] coefficients: [[ -6.17445455e+04  1.88053
371e+04  3.17237961e+02 -2.88006880e-01
-9.68870610e+03  5.75218487e+02]]
RMSE: 249311.90
```

The RMSE goes down from \$254.323,39 to \$228.024,43 with more features.

Apply learned models to predict prices of 3 houses

The first house we will use is considered an "average" house in Seattle.

In [25]:

```
house1 = sales[sales['id']==5309101200]
```


In [26]:

```
house1
```

Out[26]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft
1054	5309101200	20140605T000000	620000	4	2.25	2400	5350

1 rows × 21 columns





In [27]:

```
print (house1['price'])
```

```
1054    620000
```

```
Name: price, dtype: int64
```

In [28]:

```
house1['sqft_living']
```

Out[28]:

```
1054    2400
```

```
Name: sqft_living, dtype: int64
```

In [29]:

```
print (simple_model.predict(house1['sqft_living'].values.reshape(-1,1)))
```

```
[[ 630959.4568039]]
```


In [30]:

```
print (mult_model.predict(house1[my_features]))  
[[ 630924.33807747]]
```

In this case, the model with more features provides a worse prediction than the simpler model with only 1 feature. However, on average, the model with more features is better.

Prediction for a second, fancier house

We will now examine the predictions for a fancier house.

In [31]:

```
house2 = sales[sales['id']==1925069082]
```

In [32]:

```
house2
```

Out[32]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft
1361	1925069082	20150511T000000	2200000	5	4.25	4640	227

1 rows × 21 columns



In [33]:

```
print (house2['price'])
```

```
1361    2200000
```

```
Name: price, dtype: int64
```

In [34]:

```
print (simple_model.predict(house2['sqft_living'].values.reshape(-1,1)))
```

```
[[ 1263248.45867172]]
```

In [35]:

```
print (mult_model.predict(house2[my_features]))
```

```
[[ 1270172.16078082]]
```

In this case, the model with more features provides a better prediction. This behavior is expected here, because this house is more differentiated by features that go beyond its square feet of living space, especially the fact that it's a waterfront house.

Last house, super fancy

Our last house is a very large one owned by a famous Seattleite.

In [36]:

```
bill_gates = {'bedrooms':[8],
              'bathrooms':[25],
              'sqft_living':[50000],
              'sqft_lot':[225000],
              'floors':[4],
              'zipcode':['98039'],
              'condition':[10],
              'grade':[10],
              'waterfront':[1],
              'view':[4],
              'sqft_above':[37500],
              'sqft_basement':[12500],
              'yr_built':[1994],
              'yr_renovated':[2010],
              'lat':[47.627606],
              'long':[-122.242054],
              'sqft_living15':[5000],
              'sqft_lot15':[40000]}
```




In [37]:

```
print (simple_model.predict(pd.DataFrame(bill_gates)
['sqft_living'].values.reshape(-1,1)))
```

```
[[ 14067100.74649498]]
```

The model predicts a price of over \$14M for this house! But we expect the house to cost much more. (There are very few samples in the dataset of houses that are this fancy, so we don't expect the model to capture a perfect prediction here.)

In [38]:

```
print (mult_model.predict(pd.DataFrame(bill_gates)[my_features]))
```

```
[[ 15779944.98847018]]
```

Answers

1 - Selection and summary statistics

In [39]:

```
zip_code = sales[sales['zipcode']==98039]
```

In [40]:

```
my_features = ['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'zipcode']
```

In [41]:

```
zip_code[my_features].describe()
```

Out[41]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	zip
count	5.000000e+01	50.000000	50.000000	50.0000	50.000000	50.000000	50
mean	2.160607e+06	4.060000	3.200000	3800.9000	17403.560000	1.560000	98
std	1.166477e+06	0.890081	1.366509	1764.5025	6655.224175	0.501427	0.0
min	7.875000e+05	2.000000	1.000000	1220.0000	6572.000000	1.000000	98
25%	1.401000e+06	4.000000	2.250000	2680.0000	13797.500000	1.000000	98
50%	1.892500e+06	4.000000	3.000000	3560.0000	17188.500000	2.000000	98
75%	2.556250e+06	4.750000	3.687500	4452.5000	20052.250000	2.000000	98
max	6.885000e+06	7.000000	7.750000	9890.0000	35069.000000	2.000000	98

In [42]:

```
print ('O preço médio de venda da vizinhança é $%.2f.' % (zip_code['price'].mean()))
```

O preço médio de venda da vizinhança é \$2160606.60.

2 - Filtering data

In [43]:

```
area_filter = sales[(sales['sqft_living']>=2000) & (sales['sqft_living']<=4000)]
```

In [44]:

```
print(area_filter.head())
```

	id	date	price	bedrooms	bathrooms	sqft_living
1	6414100192	20141209T000000	538000	3	2.25	2570
10	1736800520	20150403T000000	662500	3	2.50	3560
15	9297300055	20150124T000000	650000	4	3.00	2950
21	2524049179	20140826T000000	2000000	3	2.75	3050
22	7137970340	20140703T000000	285000	5	2.50	2270

	sqft_lot	floors	waterfront	view	...	grade	sqft_above	\
1	7242	2.0	0	0	...	7	2170	
10	9796	1.0	0	0	...	8	1860	
15	5000	2.0	0	3	...	9	1980	
21	44867	1.0	0	4	...	9	2330	
22	6300	2.0	0	0	...	8	2270	

	sqft_basement	yr_built	yr_renovated	zipcode	lat	long	\
1	400	1951	1991	98125	47.7210	-122.319	
10	1700	1965	0	98007	47.6007	-122.145	
15	970	1979	0	98126	47.5714	-122.375	
21	720	1968	0	98040	47.5316	-122.233	
22	0	1995	0	98092	47.3266	-122.169	

	sqft_living15	sqft_lot15
1	1690	7639
10	2210	8925
15	2140	4000
21	4110	20336
22	2240	7005

[5 rows x 21 columns]

In [45]:

```
print ('0 percentual de casas que possuem áreas entre 2000 e 4000 pés quadrados é de:
%.2f.' % (float(area_filter['zipcode'].count())/(sales['zipcode'].count()) * 100))
```

0 percentual de casas que possuem áreas entre 2000 e 4000 pés quadrados é de: 42.66.

3 - Building a regression model with several more features

In [46]:

```
advanced_features = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'zip
code', 'condition', 'grade',
                    'waterfront', 'view', 'sqft_above', 'sqft_basement', 'yr_built']
```